

CODE CLAUSE PROJECT

PROJECT NAME - Churn Prediction in Telecom Industry using Logistic Regression

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import plotly.express as px
```

Importing Datasets

```
In [2]: telecom_data = pd.read_csv('D:/CODE CLAUSE DATA SCIENCE/CHURN PREDICTION/archive/Telc
```

```
In [3]: telecom_data.head()
```

```
Out[3]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleL
--	------------	--------	---------------	---------	------------	--------	--------------	-----------

0	7590-VHVEG	Female	0	Yes	No	1	No	No ph ser
---	------------	--------	---	-----	----	---	----	--------------

1	5575-GNVDE	Male	0	No	No	34	Yes	
---	------------	------	---	----	----	----	-----	--

2	3668-QPYBK	Male	0	No	No	2	Yes	
---	------------	------	---	----	----	---	-----	--

3	7795-CFOCW	Male	0	No	No	45	No	No ph ser
---	------------	------	---	----	----	----	----	--------------

4	9237-HQITU	Female	0	No	No	2	Yes	
---	------------	--------	---	----	----	---	-----	--

5 rows × 21 columns

```
In [4]: telecom_data.shape
```

```
Out[4]: (7043, 21)
```

```
In [5]: telecom_data.describe()
```

Out[5]:

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000


In [6]: *# Checking Null Values*
 telecom_data.notnull().sum()

Out[6]: customerID 7043
 gender 7043
 SeniorCitizen 7043
 Partner 7043
 Dependents 7043
 tenure 7043
 PhoneService 7043
 MultipleLines 7043
 InternetService 7043
 OnlineSecurity 7043
 OnlineBackup 7043
 DeviceProtection 7043
 TechSupport 7043
 StreamingTV 7043
 StreamingMovies 7043
 Contract 7043
 PaperlessBilling 7043
 PaymentMethod 7043
 MonthlyCharges 7043
 TotalCharges 7043
 Churn 7043
 dtype: int64

In [7]: *#There is no missing value in our*


In [8]: *%matplotlib inline*

In [9]: telecom_hist = px.histogram(telecom_data, x='gender',color='Churn',marginal='box', co
 telecom_hist.update_layout(bargap=0.2)





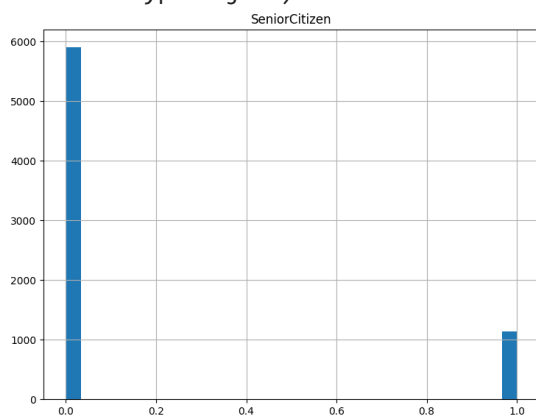
```
In [10]: plt.bar(telecom_data['gender'],telecom_data['Churn'])
```

```
Out[10]: <BarContainer object of 7043 artists>
```



```
In [11]: telecom_data.hist(bins = 30, figsize=(20,15))
```

```
Out[11]: array([[<Axes: title={'center': 'SeniorCitizen'}>,
                <Axes: title={'center': 'tenure'}>],
               [<Axes: title={'center': 'MonthlyCharges'}>, <Axes: >]],
          dtype=object)
```




```
In [12]: sns.pairplot(telecom_data)
```

C:\Users\USER\AppData\Local\Programs\Python\Python311\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning:

The figure layout has changed to tight

```
Out[12]: <seaborn.axisgrid.PairGrid at 0x15a311c6d10>
```



Cleaning Data

```
In [13]: #Removing gender, customerID, tenure they are not usefull
```


```
In [14]: col = ['gender', 'customerID', 'tenure']
telecom_data = telecom_data.drop(col, axis = 1)
```

```
In [15]: sns.pairplot(telecom_data)
```

C:\Users\USER\AppData\Local\Programs\Python\Python311\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning:

The figure layout has changed to tight

```
Out[15]: <seaborn.axisgrid.PairGrid at 0x15a31d90dd0>
```



```
In [16]: telecom_data.head()
```

```
Out[16]:
```

	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineService
0	0	Yes	No	No	No phone service	DSL	
1	0	No	No	Yes	No	DSL	
2	0	No	No	Yes	No	DSL	
3	0	No	No	No	No phone service	DSL	
4	0	No	No	Yes	No	Fiber optic	

```
In [17]: telecom_data['TotalCharges'].notnull().sum()
```

```
Out[17]: 7043
```

```
In [18]: telecom_data['MonthlyCharges'].describe()
```

```
Out[18]: count      7043.000000
         mean        64.761692
         std         30.090047
         min         18.250000
         25%         35.500000
         50%         70.350000
         75%         89.850000
         max         118.750000
         Name: MonthlyCharges, dtype: float64
```

```
In [19]: telecom_data['TotalCharges'].describe()
         #the data type of the Total Charges is Object so we will change that
```

```
Out[19]: count      7043
         unique     6531
         top
         freq        11
         Name: TotalCharges, dtype: object
```

```
In [20]: #due to string(" ") at 488 position you can not change the TotalCharges into Int
         #so we will be removing/replacing that string which is --> " "
         telecom_data['TotalCharges'] = telecom_data['TotalCharges'].replace(" ",np.nan)
         telecom_data['TotalCharges'] = pd.to_numeric(telecom_data['TotalCharges'], errors = '
         #dropping all the rows in which there is a null value
         telecom_data = telecom_data.dropna(how = "any", axis = 0) #removing all the rows whic
```

```
In [21]: telecom_data['TotalCharges'].describe()
```

```
Out[21]: count      7032.000000
         mean      2283.300441
         std       2266.771362
         min       18.800000
         25%       401.450000
         50%      1397.475000
         75%      3794.737500
         max      8684.800000
         Name: TotalCharges, dtype: float64
```

```
In [22]: telecom_data.notnull().sum()
```

```
Out[22]: SeniorCitizen      7032
         Partner            7032
         Dependents         7032
         PhoneService       7032
         MultipleLines       7032
         InternetService    7032
         OnlineSecurity     7032
         OnlineBackup       7032
         DeviceProtection   7032
         TechSupport        7032
         StreamingTV        7032
         StreamingMovies    7032
         Contract           7032
         PaperlessBilling   7032
         PaymentMethod      7032
         MonthlyCharges     7032
         TotalCharges       7032
         Churn              7032
         dtype: int64
```

```
In [23]: #Total Charges has null values in it
```

```
In [24]: telecom_data.isnull().sum()
```



```
Out[24]: SeniorCitizen      0
Partner      0
Dependents    0
PhoneService  0
MultipleLines 0
InternetService 0
OnlineSecurity 0
OnlineBackup  0
DeviceProtection 0
TechSupport   0
StreamingTV   0
StreamingMovies 0
Contract      0
PaperlessBilling 0
PaymentMethod 0
MonthlyCharges 0
TotalCharges  0
Churn         0
dtype: int64
```



EDA(Exploratory Data Analysis)

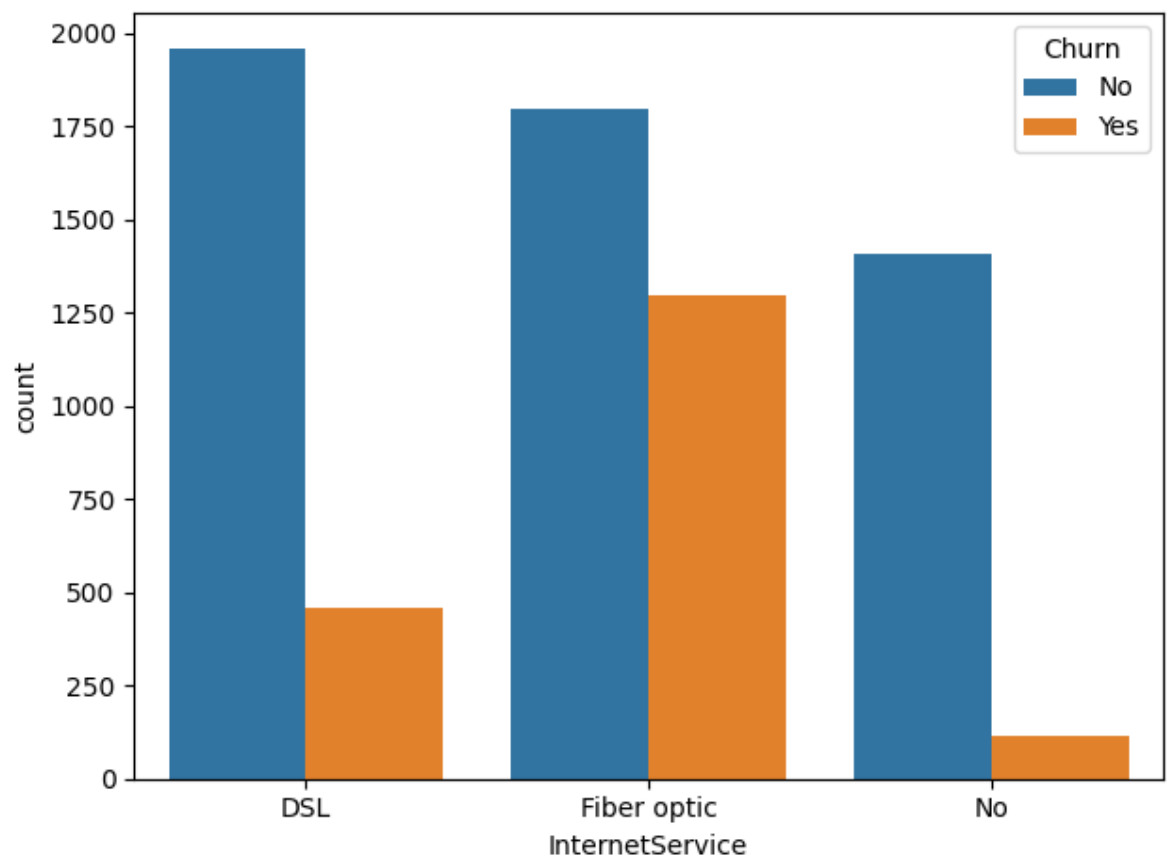

```
In [25]: telecom_data['Churn'].describe()
```



```
Out[25]: count      7032
unique         2
top           No
freq         5163
Name: Churn, dtype: object
```


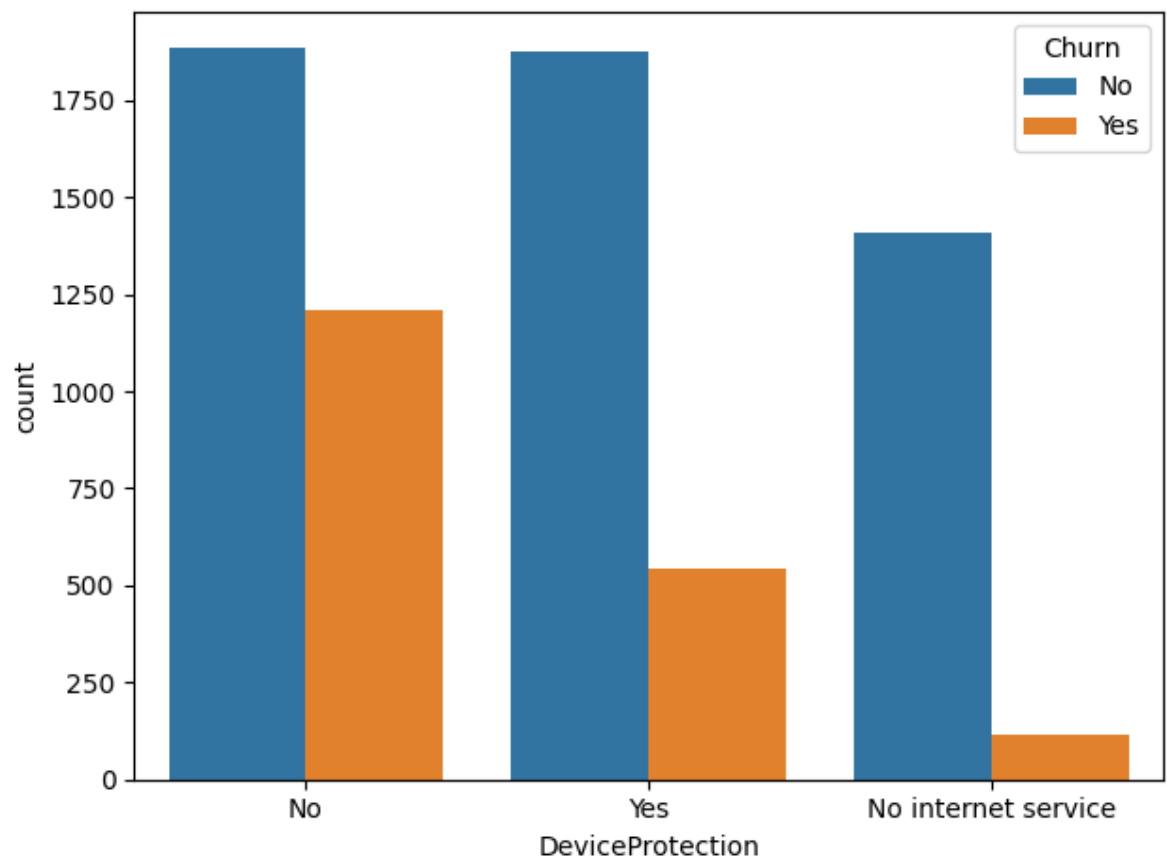
```
In [26]: for i, predictor in enumerate(telecom_data.drop(columns=['Churn', 'TotalCharges', 'Mo
ax = sns.countplot(data =telecom_data, x = predictor, hue='Churn')
if predictor == "PaymentMethod":
    ax.set_xticklabels(ax.get_xticklabels(), fontsize=7)
    plt.tight_layout()
    plt.show()
else:
    plt.tight_layout()
    plt.show()
```























```
In [27]: #converting Yes as 1 and No as 0
telecom_data["Churn"] = telecom_data["Churn"].replace(['Yes','No'],[1,0])
```

C:\Users\USER\AppData\Local\Temp\ipykernel_8512\1073425660.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
In [28]: telecom_data
```


Out[28]:

	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService	Online
0	0	Yes	No	No	No phone service	DSL	
1	0	No	No	Yes	No	DSL	
2	0	No	No	Yes	No	DSL	
3	0	No	No	No	No phone service	DSL	
4	0	No	No	Yes	No	Fiber optic	
...
7038	0	Yes	Yes	Yes	Yes	DSL	
7039	0	Yes	Yes	Yes	Yes	Fiber optic	
7040	0	Yes	Yes	No	No phone service	DSL	
7041	1	Yes	No	Yes	Yes	Fiber optic	
7042	0	No	No	Yes	No	Fiber optic	


7032 rows × 18 columns

In [29]: `telecom_data_dummies = pd.get_dummies(telecom_data)`In [30]: `telecom_data_dummies`

Out[30]:

	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	Partner_No	Partner_Yes	Depend
0	0	29.85	29.85	0	False	True	
1	0	56.95	1889.50	0	True	False	
2	0	53.85	108.15	1	True	False	
3	0	42.30	1840.75	0	True	False	
4	0	70.70	151.65	1	True	False	
...	
7038	0	84.80	1990.50	0	False	True	
7039	0	103.20	7362.90	0	False	True	
7040	0	29.60	346.45	0	False	True	
7041	1	74.40	306.60	1	False	True	
7042	0	105.65	6844.50	0	True	False	


7032 rows x 43 columns



```
In [31]: churn_corr_matrix = telecom_data_dummies.corr()
```

```
In [32]: churn_corr_matrix['Churn'].sort_values(ascending = False).plot(kind='bar',figsize = (
```

Out[32]: <Axes: >



```
In [33]: churn_corr_matrix['Churn'].sort_values(ascending = False)
```

```

Out[33]: Churn 1.000000
Contract_Month-to-month 0.404565
OnlineSecurity_No 0.342235
TechSupport_No 0.336877
InternetService_Fiber optic 0.307463
PaymentMethod_Electronic check 0.301455
OnlineBackup_No 0.267595
DeviceProtection_No 0.252056
MonthlyCharges 0.192858
PaperlessBilling_Yes 0.191454
Dependents_No 0.163128
SeniorCitizen 0.150541
Partner_No 0.149982
StreamingMovies_No 0.130920
StreamingTV_No 0.128435
StreamingTV_Yes 0.063254
StreamingMovies_Yes 0.060860
MultipleLines_Yes 0.040033
PhoneService_Yes 0.011691
PhoneService_No -0.011691
MultipleLines_No phone service -0.011691
MultipleLines_No -0.032654
DeviceProtection_Yes -0.066193
OnlineBackup_Yes -0.082307
PaymentMethod_Mailed check -0.090773
PaymentMethod_Bank transfer (automatic) -0.118136
InternetService_DSL -0.124141
PaymentMethod_Credit card (automatic) -0.134687
Partner_Yes -0.149982
Dependents_Yes -0.163128
TechSupport_Yes -0.164716
OnlineSecurity_Yes -0.171270
Contract_One year -0.178225
PaperlessBilling_No -0.191454
TotalCharges -0.199484
StreamingTV_No internet service -0.227578
OnlineSecurity_No internet service -0.227578
InternetService_No -0.227578
StreamingMovies_No internet service -0.227578
OnlineBackup_No internet service -0.227578
TechSupport_No internet service -0.227578
DeviceProtection_No internet service -0.227578
Contract_Two year -0.301552
Name: Churn, dtype: float64

```


```
In [34]: x = telecom_data_dummies.drop('Churn',axis = 1)
```

```
In [35]: x
```

Out[35]:

	SeniorCitizen	MonthlyCharges	TotalCharges	Partner_No	Partner_Yes	Dependents_No
0	0	29.85	29.85	False	True	True
1	0	56.95	1889.50	True	False	True
2	0	53.85	108.15	True	False	True
3	0	42.30	1840.75	True	False	True
4	0	70.70	151.65	True	False	True
...
7038	0	84.80	1990.50	False	True	False
7039	0	103.20	7362.90	False	True	False
7040	0	29.60	346.45	False	True	False
7041	1	74.40	306.60	False	True	True
7042	0	105.65	6844.50	True	False	True

7032 rows × 42 columns



In [36]: `y = telecom_data_dummies['Churn']`

In [37]: `y`

Out[37]:

0	0
1	0
2	1
3	0
4	1
...	...
7038	0
7039	0
7040	0
7041	1
7042	0

Name: Churn, Length: 7032, dtype: int64

In [38]: `x.shape`

Out[38]: (7032, 42)

In [39]: `y.shape`

Out[39]: (7032,)

In [40]: `y.value_counts()`

Out[40]:

Churn	
0	5163
1	1869

Name: count, dtype: int64

Variable Imbalancing

SMOTE for Imbalanced Classification with Python

```
In [41]: from imblearn.over_sampling import SMOTE
```

```
In [42]: smote = SMOTE(random_state=0)
```

```
In [43]: x_resampled_smote, y_resampled_smote = smote.fit_resample(x,y)
```

```
In [44]: y_resampled_smote.value_counts()
```

```
Out[44]: Churn
0      5163
1      5163
Name: count, dtype: int64
```

```
In [45]: x_resampled_smote
```

```
Out[45]:
```

	SeniorCitizen	MonthlyCharges	TotalCharges	Partner_No	Partner_Yes	Dependents_No
0	0	29.850000	29.850000	False	True	True
1	0	56.950000	1889.500000	True	False	True
2	0	53.850000	108.150000	True	False	True
3	0	42.300000	1840.750000	True	False	True
4	0	70.700000	151.650000	True	False	True
...
10321	0	103.976753	242.804921	False	True	True
10322	0	35.824447	35.824447	True	False	True
10323	0	44.493077	1061.960339	True	True	True
10324	0	19.363055	19.363055	True	False	True
10325	0	96.922890	96.922890	True	False	True

10326 rows × 42 columns

```
In [46]: y_resampled_smote.notnull().sum()
```

```
Out[46]: 10326
```

```
In [47]: x_resampled_smote.notnull().sum()
```

```
Out[47]: SeniorCitizen      10326
MonthlyCharges      10326
TotalCharges        10326
Partner_No          10326
Partner_Yes         10326
Dependents_No       10326
Dependents_Yes      10326
PhoneService_No     10326
PhoneService_Yes    10326
MultipleLines_No     10326
MultipleLines_No phone service 10326
MultipleLines_Yes   10326
InternetService_DSL  10326
InternetService_Fiber optic 10326
InternetService_No  10326
OnlineSecurity_No   10326
OnlineSecurity_No internet service 10326
OnlineSecurity_Yes  10326
OnlineBackup_No     10326
OnlineBackup_No internet service 10326
OnlineBackup_Yes    10326
DeviceProtection_No 10326
DeviceProtection_No internet service 10326
DeviceProtection_Yes 10326
TechSupport_No      10326
TechSupport_No internet service 10326
TechSupport_Yes     10326
StreamingTV_No      10326
StreamingTV_No internet service 10326
StreamingTV_Yes     10326
StreamingMovies_No  10326
StreamingMovies_No internet service 10326
StreamingMovies_Yes 10326
Contract_Month-to-month 10326
Contract_One year   10326
Contract_Two year   10326
PaperlessBilling_No 10326
PaperlessBilling_Yes 10326
PaymentMethod_Bank transfer (automatic) 10326
PaymentMethod_Credit card (automatic) 10326
PaymentMethod_Electronic check 10326
PaymentMethod_Mailed check 10326
dtype: int64
```

```
In [48]: from sklearn.linear_model import LogisticRegression
```

```
In [49]: #checking on imbalance data
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=42)
```

```
In [50]: LogReg = LogisticRegression(solver='lbfgs', max_iter=1000,multi_class='multinomial')
```

```
In [51]: LogReg.fit(x_train,y_train)
```

```
Out[51]: LogisticRegression
LogisticRegression(max_iter=1000, multi_class='multinomial')
```

```
In [52]: y_pred = LogReg.predict(x_test)
```

```
In [53]: from sklearn.metrics import accuracy_score
```

```
In [54]: accuracy_score(y_test,y_pred)
```

```
Out[54]: 0.7853589196872779
```

```
In [55]: #checking on balanced data  
x_smote_train,x_smote_test,y_smote_train,y_smote_test = train_test_split(x_resampled_
```

```
In [56]: LogReg.fit(x_smote_train,y_smote_train)
```

```
Out[56]: ▼ LogisticRegression  
LogisticRegression(max_iter=1000, multi_class='multinomial')
```

```
In [57]: y_smote_pred = LogReg.predict(x_smote_test)
```

```
In [58]: accuracy_score(y_smote_test,y_smote_pred)
```

```
Out[58]: 0.8388189738625363
```

```
In [59]: from sklearn.preprocessing import StandardScaler
```

```
In [60]: std = StandardScaler()
```

```
In [61]: std_train = std.fit_transform(x_smote_train)  
std_test = std.transform(x_smote_test)
```

```
In [62]: LogReg.fit(std_train,y_smote_train)
```

```
Out[62]: ▼ LogisticRegression  
LogisticRegression(max_iter=1000, multi_class='multinomial')
```

```
In [63]: std_pred = LogReg.predict(std_test)
```

```
In [64]: accuracy_score(std_pred,y_smote_test)
```

```
Out[64]: 0.8407550822846079
```

```
In [65]: np.where(std_pred!=y_smote_test)
```



```
Out[65]: (array([ 14,  20,  24,  31,  43,  48,  49,  57,  60,  63,  80,
                  81,  83,  87,  90,  98, 100, 102, 107, 108, 117, 118,
                  125, 126, 130, 136, 161, 162, 183, 193, 194, 207, 230,
                  236, 272, 274, 287, 289, 291, 296, 300, 306, 313, 321,
                  327, 328, 329, 330, 333, 335, 341, 346, 348, 359, 376,
                  380, 393, 397, 400, 414, 415, 421, 425, 427, 428, 434,
                  435, 439, 442, 449, 451, 463, 479, 489, 490, 491, 499,
                  509, 515, 521, 530, 532, 543, 546, 551, 555, 556, 562,
                  563, 571, 573, 575, 585, 588, 595, 602, 608, 612, 625,
                  629, 637, 645, 661, 691, 695, 705, 710, 724, 734, 739,
                  756, 757, 760, 774, 777, 783, 785, 789, 790, 791, 794,
                  799, 805, 814, 820, 821, 841, 855, 862, 865, 866, 869,
                  870, 874, 883, 888, 899, 902, 904, 909, 912, 921, 927,
                  929, 932, 938, 940, 947, 951, 954, 962, 964, 967, 970,
                  973, 974, 986, 1003, 1005, 1015, 1037, 1043, 1045, 1046, 1047,
                  1052, 1064, 1066, 1075, 1076, 1095, 1111, 1112, 1115, 1126, 1131,
                  1134, 1135, 1137, 1141, 1146, 1148, 1152, 1154, 1157, 1159, 1165,
                  1167, 1170, 1181, 1187, 1190, 1198, 1205, 1212, 1216, 1217, 1225,
                  1236, 1239, 1241, 1242, 1243, 1250, 1253, 1265, 1273, 1282, 1287,
                  1292, 1300, 1301, 1309, 1316, 1332, 1350, 1362, 1375, 1377, 1390,
                  1392, 1407, 1413, 1414, 1428, 1439, 1440, 1457, 1458, 1460, 1463,
                  1470, 1476, 1483, 1484, 1491, 1496, 1504, 1507, 1514, 1517, 1518,
                  1525, 1539, 1547, 1549, 1553, 1562, 1575, 1579, 1580, 1584, 1596,
                  1605, 1606, 1614, 1617, 1622, 1625, 1629, 1630, 1631, 1647, 1651,
                  1654, 1666, 1667, 1678, 1681, 1683, 1693, 1719, 1721, 1724, 1732,
                  1733, 1741, 1748, 1752, 1754, 1768, 1777, 1787, 1791, 1798, 1801,
                  1815, 1818, 1822, 1827, 1831, 1832, 1834, 1835, 1836, 1837, 1843,
                  1850, 1854, 1874, 1880, 1890, 1902, 1909, 1912, 1917, 1919, 1920,
                  1930, 1940, 1944, 1953, 1965, 1979, 1980, 1984, 1997, 2002, 2012,
                  2018, 2028, 2029, 2038, 2039, 2041, 2043, 2046, 2050, 2056],
          dtype=int64),)
```

```
In [66]: y_smote_test.shape
```

```
Out[66]: (2066,)
```

```
In [ ]:
```